


Toward Discovery of the Artist's Style

© GHEENT, KATHEDRALE KERKFABRIEK. PHOTO COURTESY OF KIKIRPA, BRUSSELS. MAGNIFYING GLASS—IMAGE LICENSED BY INGRAM PUBLISHING.



Signal Processing for Art Investigation

[Learning to recognize artists by their artworks]

Author attribution through the recognition of visual characteristics is a commonly used approach by art experts. By studying a vast number of artworks, art experts acquire the ability to recognize the unique characteristics of artists. In this article, we present an approach that uses the same principles to discover the characteristic features that determine an artist's touch. By training a convolutional neural network (PigeoNET) on a large collection of digitized artworks to perform the task of automatic artist attribution, the network is encouraged to discover artist-specific visual features. The trained network is shown to be capable of attributing previously

unseen artworks to the actual artists with an accuracy of more than 70%. In addition, the trained network provides fine-grained information about the artist-specific characteristics of spatial regions within the artworks. We demonstrate this ability by means of a single artwork that combines characteristics of two closely collaborating artists. PigeoNET generates a visualization that indicates for each location on the artwork who is the most likely artist to have contributed to the visual characteristics at that location. We conclude that PigeoNET represents a fruitful approach for the future of computer-supported examination of artworks.

INTRODUCTION

Identifying the artist of an artwork is a crucial step in establishing its value from a cultural, historical, and economic perspective. Typically, the attribution is performed by an experienced art expert

with a long-standing reputation and an extensive knowledge of the features characteristic of the alleged artist and contemporaries.

Art experts acquire their knowledge by studying a vast number of artworks accompanied by descriptions of the relevant characteristics (features) [1]. For instance, the characteristic features of Vincent van Gogh during his later French period include the outlines painted around objects, complementary colors [2], and rhythmic brush strokes [3]. As Van Dantzig [4] claimed in the context of his *Pictology* approach, describing works by an artist in terms of visual features enables the attribution of works to artists (see also [5]).

The advent of computers and high-resolution digital reproductions of artworks gave rise to attempts to partially automate the attribution of artworks [6]–[8]. Given the appropriate visual features, machine-learning algorithms may automatically attribute artworks to their artists. As was (and still is) common practice in traditional machine learning, feature engineering, i.e., finding or defining the appropriate features, is critical to the success of the automatic attribution task. Close cooperation with art historians and conservators facilitated the feature engineering for artist attribution, which led to promising results in the automatic attribution of artworks by van Gogh and his contemporaries [3], [6], [9], [10], highlighting the value of automatic approaches as a tool for art experts.

Despite the success of feature engineering, these early attempts were hampered by the difficulty to acquire explicit knowledge about all the features associated with the artists of artworks. Understandably, the explicit identification of characteristic features posed a challenge to art experts, because (as is true for most experts) their expertise is based on tacit knowledge that is difficult to verbalize [11]. By adopting a method capable of automatically recognizing the characteristics that are known to be important for the task at hand, the tacit knowledge of art experts may be operationalized [12].

Feature learning is an alternative to feature engineering that learns features directly from the data [12]. Feature learning is much more data intensive than feature engineering, because it requires a large number of examples to discover the characteristic features. In recent years, feature learning has shown great promise by taking advantage of deep architectures, machine learning methods inspired by biological neural networks. A typical example of a deep architecture is a convolutional neural network, which, when combined with a powerful learning algorithm, is capable of discovering (visual) features. Convolutional neural networks outperform all existing learning algorithms on a variety of very challenging image classification tasks [13]. To our knowledge, convolutional neural networks have not yet been applied for automated artist attribution. The objective of this article is to present a novel and transparent way of performing automatic artist attribution of artworks by means of convolutional neural networks.

IDENTIFYING THE ARTIST OF AN ARTWORK IS A CRUCIAL STEP IN ESTABLISHING ITS VALUE FROM A CULTURAL, HISTORICAL, AND ECONOMIC PERSPECTIVE.

When using only visual information, the following question may be raised: Is automatic artist attribution possible at all? It has been frequently argued by scholars working in the art domain that semantic or historical knowledge, as well as technical and analytical information are pivotal in the attribution of artworks. The feasibility of image-based automatic

artist attribution is supported by biological studies. Pigeons [14] and honeybees [15] can be successfully trained to discriminate between artists, with pigeons correctly attributing an art work in 90% of the cases in a binary Monet–Picasso attribution task. This shows that a visual

system without higher cognitive functions is capable of learning the visual characteristics present in artworks. While it is unlikely that a perfect result can be achieved without incorporating additional information, these findings do pave the way for an attribution approach that learns to recognize visual features from data rather than from prior knowledge.

In this article, we present PigeoNET, a convolutional neural network corresponding to the AlexNET architecture described in [13] to which we added a visualization component due to [16]. PigeoNET is applied to an artist attribution task by training it on artworks. As such, PigeoNET performs a task similar to the pigeons in [14], by performing artist attribution based solely on visual characteristics. This implies that, in addition to authorship, PigeoNET may also take visual characteristics into consideration that relate indirectly to the artist (e.g., the choice of materials or tools used by the artist) or that are completely unrelated to the artist (e.g., reproduction characteristics such as lighting and digitization procedure). To ensure that the visual characteristics on which the task is solved by PigeoNET make sense, human experts are needed to assess the relevance of the acquired mapping from images of artworks to artists. Our visualization method allows for the visual assessment by experts of the characteristic regions of artworks.

In our artist attribution experiments, we consider three sources of variation in the training set and assess their effects on attribution performance: 1) heterogeneity versus homogeneity of classes (types of artworks), 2) number of artists, and 3) number of artworks per artist.

After training, the performance of PigeoNET will be assessed in two ways: 1) by determining how well it attributes previously unseen artworks and 2) by generating visualizations that reveal artwork regions characteristic of the artist, or in case of artworks that are likely created by two or more artists, generating visualizations that reveal which regions belong to which artist, and could aid in answering outstanding art historical questions.

PigeoNET

A convolutional neural network can learn to recognize the visually characteristic features of an artist by adapting filters to respond to the presence of these features in an image [17]. The filters are adapted to respond to a feature by adjusting the parameters, or weights, of the filters until a suitable configuration is found. The

proper weights for this configuration are obtained by means of a learning algorithm called *back-propagation* [18], which requires no prior knowledge other than the input images and a label (e.g., the artist who created it). In the case of artist attribution, the network will learn to recognize features that are regarded as characteristic of a certain artist, allowing us to discover these characteristics. PigeoNET is a convolutional neural network designed to learn the characteristics of artists and their artworks, so as to recognize and identify their authorship.

The filters in a convolutional neural network are grouped into layers, where the first layer is directly applied to images, and subsequent layers to the responses generated by previous layers. By stacking layers to create a multilayer architecture the filters can respond to increasingly complex features with each subsequent layer. The filters in the initial layers respond to low-level visual patterns, akin to Gabor filters [19], whereas the final layers of filters respond to visual characteristic features specific to artists.

Because convolution is used to apply the filters to an image, or the response of a previous layer, the layers of filters are referred to as *convolutional layers*. The advantage of a convolutional layer, over a traditional neural network layer, is that the weights are shared, allowing the adaptive filters to respond to characteristic features irrespective of their position or location in the input [18]. To learn a mapping from the filter responses to a certain artist the convolutional layers are, typically, followed by a number of fully connected layers that translate the presence and intensity of the filter responses to a single certainty score per artist. The certainty score for an artist is high whenever the responses for filters corresponding to that artist are strong, conversely, the certainty score is low when the filter responses are weak or nonexistent. Thus, an unseen artwork can be attributed to an artist for whom the certainty score is the highest.

PIGEONET IS A CONVOLUTIONAL NEURAL NETWORK DESIGNED TO LEARN THE CHARACTERISTICS OF ARTISTS AND THEIR ARTWORKS, SO AS TO RECOGNIZE AND IDENTIFY THEIR AUTHORSHIP.

VISUALIZATION OF ARTIST-CHARACTERISTIC REGIONS

While PigeoNET's attribution of an artwork is based on the entire artwork, regions containing visual elements characteristic for an artist are assigned more weight than others to achieve a correct attribution [20]. To increase our understanding of the attribution performed by PigeoNET, we aim to visualize such artist-characteristic regions. Several methods have been proposed for visualizing trained convolutional neural networks [16], [21] and other layered algorithms [22]. We adopt the occlusion sensitivity testing method proposed by [16] for obtaining visualizations of artist-characteristic regions, which can be considered a weakly supervised

localization method. By systematically occluding a small image region of an artwork, the importance of the occluded region is determined by observing the change in the certainty score for the correct artist. When an occluded region is very important (or highly characteristic) for correctly identifying the artist, there will be a significant drop in the certainty score generated by PigeoNET. Inversely, occluding a region that is atypical for the correct artist will result in an increase in the certainty score. A region for which occlusion results in a drop of the certainty score is considered characteristic for the artist under consideration. This approach to creating visualizations allows us to show the approximate areas of an artwork which are representative of an artist.

As an illustration, Figure 1 depicts *The Feast of Achelouïs* by two artists; Peter Paul Rubens painted the persons and Jan Brueghel the Elder painted the scenery [23]. Although there is no single correct artist, the certainty score for Brueghel would decrease if the scenery were to be occluded, whereas the certainty score for Rubens would drop if the figures were occluded. Even when only part of the figures or part of the scenery were to be occluded, we see a drop in confidence scores. In a similar vein, when even smaller regions of the painting have been occluded, it becomes possible to identify important regions on a much more detailed scale.

AUTHOR ATTRIBUTION EXPERIMENT

The goal of an artist attribution task is to attribute an unseen artwork to the artist who created it. To be able to perform this task adequately, PigeoNET needs to discover features that distinguish an artist from other artists but especially to discover features that are characteristic of each artist. In the rest of this section we will discuss the data set, network architecture, training procedure, evaluation procedure, and the results.

EXPERIMENTAL SETUP

DATA SET

The characteristic features of an artist can be discovered by studying artworks that are representative of that artist. Yet,



[FIG1] Peter Paul Rubens and Jan Brueghel the Elder's work: *The Feast of Achelouïs* (45.141) in The Metropolitan Museum of Art's *Heilbrunn Timeline of Art History* (2000–present). (Image used courtesy of The Metropolitan Museum of Art; <http://www.metmuseum.org/toah/works-of-art/45.141>.)

obtaining a sufficiently large sample of such images is problematic, given the lack of (automatic) methods and criteria to determine whether an artwork is representative. A commonly utilized approach to circumvent the need for a representative sample is to take a very large sample. As such, a data set that contains a large number of images, and a large number of images per artist, is required.

The Rijksmuseum Challenge data set [24] consists of 112,039 digital photographic reproductions of artworks by 6,629 artists exhibited in the Rijksmuseum in Amsterdam, The Netherlands. All artworks were digitized under controlled settings. Within the set there are 1,824 different types of artworks and 406 annotated materials, such as paper, canvas, porcelain, iron, and wood. To our knowledge, this is the largest available image data set of artworks and the only data set that meets our requirements.

We divided the Rijksmuseum Challenge data set into a training, validation, and test set (cf. [24]). In this article, these sets are used to train PigeoNET, to optimize the hyperparameters, and to evaluate the performance of PigeoNET on unseen examples, respectively. The data set contains a number of artworks that lack a clear attribution, these are labeled as either “Anonymous” or “Unknown.” We chose to exclude these artworks because our objective is to relate visual features to specific artists.

While the Rijksmuseum Challenge data set contains a large number of images of many different types of artworks by a large number of artists, there are many artists for whom only a few artworks are available or artists who have created many different types of artworks. As stated earlier, these variations might influence the performance of PigeoNET in nonobvious ways. To this end, we consider the following three sources of variation: 1) heterogeneity versus

[TABLE 1] A LIST OF THE 34 ARTISTS WITH AT LEAST 256 ART-WORKS AND THE DISTRIBUTION OF ARTWORKS OVER MAIN TYPES (PRINTS, DRAWINGS, AND OTHER).

NUMBER	NAME	PRINTS	DRAWINGS	OTHER
1	HEINRICH ALDEGREVER	347	27	—
2	ERNST WILLEM JAN BAGELAAR	400	27	—
3	BOËTIUS ADAMSZ. BOLSWERT	592	—	—
4	SCHELTE ADAMSZ. BOLSWERT	398	—	—
5	ANTHONIE VAN DEN BOS	531	3	—
6	NICOLAES DE BRUYN	515	2	—
7	JACQUES CALLOT	1,008	4	1
8	ADRIAEN COLLAERT	648	1	—
9	ALBRECHT DÜRER	480	9	2
10	SIMON FOKKE	1,177	90	—
11	JACOB FOLKEMA	437	4	3
12	SIMON FRISIUS	396	—	—
13	CORNELIS GALLE (I)	421	—	—
14	PHILIPS GALLE	838	—	—
15	JACOB DE GHEYN (II)	808	75	10
16	HENDRICK GOLTZIUS	763	43	4
17	FRANS HOGENBERG	636	—	4
18	ROMEYN DE HOOGHE	1,109	5	5
19	JACOB HOUBRAKEN	1,105	42	1
20	PIETER DE JODE (II)	409	1	—
21	JEAN LEPAUTRE	559	—	1
22	CASPAR LUYKEN	359	18	—
23	JAN LUYKEN	1,895	33	—
24	JACOB ERNST MARCUS	372	23	2
25	JACOB MATHAM	546	4	—
26	MEISSENER PORZELLAN MANUFAKTUR	—	—	1,003
27	PIETER NOLPE	344	2	—
28	CRISPIJN VAN DE PASSE (I)	841	15	—
29	JAN CASPAR PHILIPS	401	17	—
30	BERNARD PICART	1,369	132	3
31	MARCANTONIO RAIMONDI	448	2	—
32	REMBRANDT HARMENSZ. VAN RIJN	1,236	119	29
33	JOHANN SADELER (I)	578	1	—
34	REINIER VINKELES	573	50	—

homogeneity of classes (types of artworks), 2) number of artists, and 3) number of artworks per artist.

Two main types of subsets were defined to assess the effect of heterogeneity versus homogeneity of artworks: type A (for “All”) and

[TABLE 2] AN OVERVIEW OF SUBSETS AND THE NUMBER OF TRAINING, VALIDATION, AND TEST IMAGES PER SUBSET.

SUBSETS	NUMBER OF EXAMPLES PER ARTIST	NUMBER OF ARTISTS (CLASSES)	NUMBER OF TRAINING IMAGES	NUMBER OF VALIDATION IMAGES	NUMBER OF TEST IMAGES
A	10	958	56,024	7,915	15,860
	64	197	37,549	5,323	10,699
	128	97	28,336	4,063	8,058
	256	34	17,029	2,489	4,838
P1	10	673	44,539	6,259	12,613
	64	165	31,655	4,484	8,983
	128	78	23,750	3,408	6,761
	256	29	14,734	2,171	4,200
P2	128	26	3,328	1,209	2,277
	128	39	4,992	1,521	2,970
	128	52	6,656	2,160	4,341
	128	78	9,984	3,408	6,761
P3	10	78	780	3,408	6,761
	64	78	4,992	3,408	6,761
	128	78	9,984	3,408	6,761

The subsets are labeled by their types. Type A (“All”) are subsets containing varying artworks, examples, and examples per artist. Type P (“Prints”) refers to subsets of prints only. P1: varying numbers of artworks, examples, and examples per artist. P2: number of examples constant (128). P3: number of artists constant (78). For A and P1, the numbers of examples per artists represent the minimum numbers, while for P2 and P3, these numbers represent the exact number of artworks per artist.

[TABLE 3] THE MCA FOR THE ARTIST ATTRIBUTION TASK ON THE 15 DATA SUBSETS.

SUBSETS	NUMBER OF EXAMPLES PER ARTIST	NUMBER OF ARTISTS (CLASSES)	MCA
A	10	958	52.5
	64	197	68.2
	128	97	74.5
	256	34	78.3
P1	10	673	60.0
	64	165	70.2
	128	78	73.3
	256	29	78.8
P2	128	26	63.9
	128	39	55.6
	128	52	52.7
	128	78	52.0
P3	10	78	13.1
	64	78	38.0
	128	78	52.0

Bold values indicate the best result per type; the overall best result is underlined.

and type P (for “Prints”), respectively. As is evident from Table 1, prints form the majority of artworks in the Rijksmuseum Challenge data set. The homogeneous type of subsets (P) has three forms: P1, P2, and P3. Subsets of type P1 have varying numbers of artists and artworks per artist (as is the case for A). Subsets of type P2 have a fixed number of artworks per artist. Finally, subsets of type P3 have a fixed number of artists. We remark that the number of examples per artist for the subsets in A and P1 are minimum values. For very productive artists these subsets may include more artworks. For subsets of types P2 and P3, the number of examples is exact and constitutes a random sample of the available works per artist. A detailed overview of the resulting 15 subsets is listed in

Table 2. For the heterogeneous subset of at least 256 artworks of type A, Table 2 provides a more detailed listing that specifies the three most prominent categories: “Prints,” “Drawings,” and “Other.” (The largest subsets for P2 and P3 are identical, but are reported twice for clarity.) The “Other” category includes a variety of different artwork types, including 35 paintings.

All images were down-sampled to 256×256 pixels following the procedure described in [13], to adhere to the fixed input size requirement of the network architecture, and are normalized at runtime by subtracting the mean image as calculated on the training set.

ARCHITECTURE

The architecture of PigeoNET is based on the Caffe [25] implementation (see http://caffe.berkeleyvision.org/model_zoo.html) of the network described in [13], and consists of five convolutional layers and three fully connected layers. The number of output nodes of the last fully connected layer is equal to the number of artists in the data set, ranging from 958 to 26 artists.

TRAINING

An effective training procedure was used (cf. [13]), in that the learning rate, momentum, and weight decay hyperparameters were assigned the values of 10^{-2} , 0.9, and $5 \cdot 10^{-4}$. The learning rate was decreased by a factor ten whenever the error on the validation set stopped decreasing. The data augmentation procedure consisted of random crops and horizontal reflections. While orientation is an important feature, to detect authorship the horizontal reflections were used to create a larger sample size, as it effectively doubles the amount of available training data, providing PigeoNET with sufficient data from which to learn, while possibly negatively impacting PigeoNET’s ability to pick up on orientation clues to perform classification. In contrast to [13], only a single crop per image was used during training, with crops of size 227×227 pixels, and the batch size was set to 256 images per batch.

All training was performed using the Caffe framework [25] on a NVIDIA Tesla K20m card and took between several hours and several days, depending on the size of the subset.

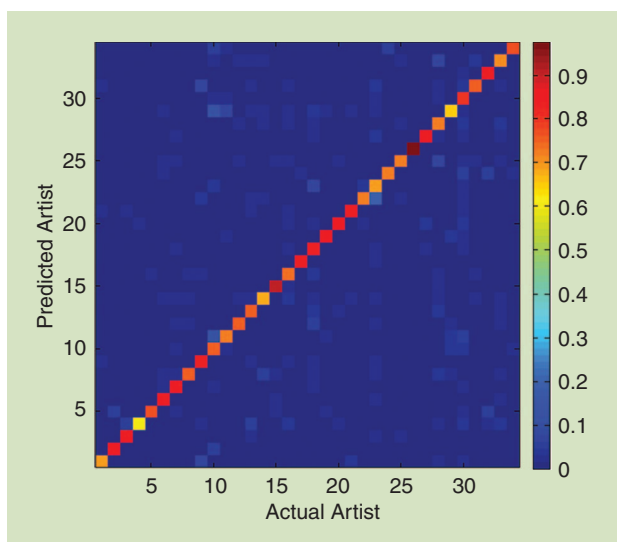
EVALUATION

The objective of the artist attribution task is to identify the correct artist for each unseen artwork in the test set. To this end the performance is measured using the mean class accuracy (MCA), which is the average of the accuracies for all artists. This makes sure that the overall performance is not heavily biased by the performance on a single artist.

During testing the final prediction is averaged over the output of the final softmax layer of the network for ten crops per image. These crops are the four corner patches and the central patch plus their horizontal reflections.

RESULTS

The results of the artist attribution task are listed in Table 3. The results on the artist attribution task show that the three sources of



[FIG2] A confusion matrix for all artists with at least 256 training examples of all artwork types. The rows represent the artist estimates and the columns the actual artists. Row and column numbers (from left to right and from bottom to top) correspond to those as listed in Table 2.

variation, [heterogeneity versus homogeneity of classes (types of artworks), number of artists, and number of artworks per artist] affect the performance in different ways. The effect of heterogeneity versus homogeneity can be assessed by comparing the results for A and P1. The results obtained with P1 are slightly better than those obtained with A (except for 128 examples per artist). However, A and P1 differ also in number of artists, which is likely to affect the performances as is evident from the results on P2 and P3.

The total number of artists (P2) and the number of examples per artist (P3) have a more prominent effect on the attribution performance of PigeoNET. Increasing the number of artists while keeping the number of examples per artist constant (as in P2) leads to a decrease in performance. With more examples per artist (P3) the performance increases tremendously.

Our results suggest that the effects of the number of artists and the number of examples per artist are closely related. This agrees with the findings reported in [13] and leads to the observation that by considering more examples per artist the number of artists to be modeled can be increased.

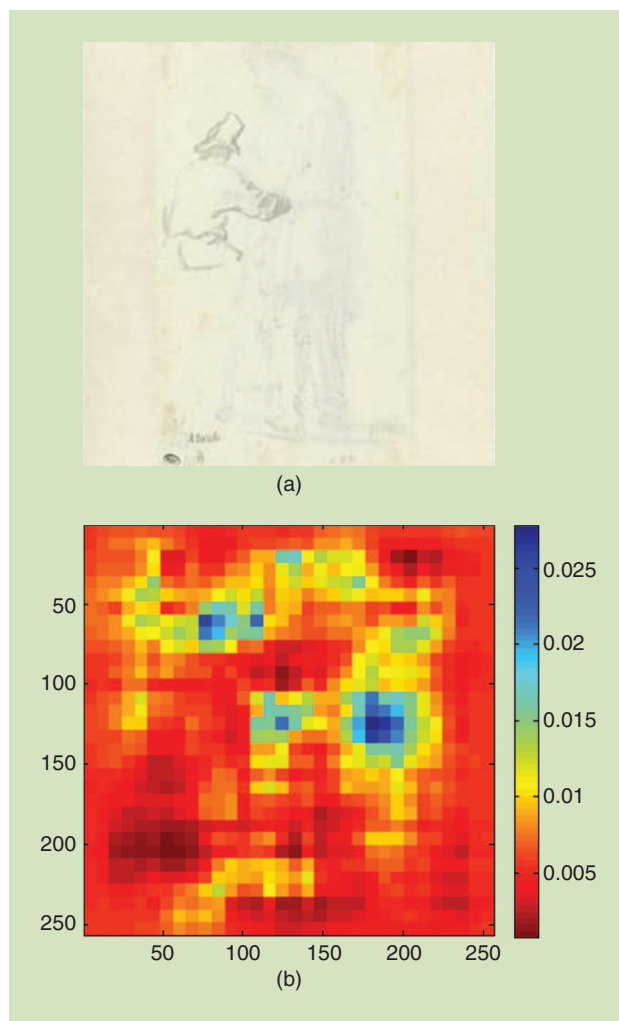
The subsets of type A are comparable to the subsets used in [24], who obtain a comparable MCA of 76.3 on a data set containing 100 artists using SIFT (which stands for “scale-invariant feature transform”) features, Fisher vectors, and one-versus-rest classification.

Figure 2 shows a visualization of the confusion matrix for the subset with at least 256 examples of all artwork types. The rows and columns correspond to the artists in Table 2. The rows represent the artist estimates by PigeoNET, the columns the actual artists. The diagonal entries represent correct attributions, which are color coded.

Upon further analysis of the results for the 256 example subset (A) of all artwork types it can be observed that the best artist-specific classification accuracy (97.5%) is obtained for Meissener Porzellan Manufaktur, a German porcelain manufacturer (class 26). Among the different types of artworks in the data set, these porcelain artworks are visually the most distinctive. Given that the visual characteristics of porcelain differ considerably from all other artworks in the data set, it is not surprising that the highest classification accuracy is achieved for this class.

The worst artist-specific classification accuracy (60.6%) is achieved for Schelte Bolswert (class 4), as indicated by the yellow square on the diagonal in the confusion matrix (fourth row from below, fourth column from left). The low accuracy may be partially explained by the confusion between Schelte Bolswert and his older brother and instructor Boëtius Bolswert (class 3). Yet, because the classification accuracy for Boëtius Bolswert (86.3) seems much less affected by the confusion, an alternative possibility is that PigeoNET is more inclined to assign visual characteristics that are present in their works to Boëtius Bolswert because his works appear more frequently in the data set.

In a similar vein, the misclassifications that occur between Fokke Simon (10) and Jan Caspar Philips (29), and between Jan Luyken (23) and Caspar Luyken (22), are notable. Fokke Simon was a student of Jan Caspar Philips, and Jan and Caspar



[FIG3] The (a) image and (b) heat map of *Studie van een op de rug geziene man* (1629–1630) by Rembrandt Harmenszoon van Rijn. The lower (red) values in the heat map correspond to greater importance in correctly identifying Rembrandt.

Luyken were father and son. Both pairs of artists have worked together on several artworks in the Rijksmuseum Challenge data set, despite the label in the data set indicating that these artworks belong to only one of these artists. We became aware of these potential dual-authorship cases after having performed our main experiment. Dual-authorship cases will be examined in more detail through visualizations in the section “Deciding Between Two Artists.”

VISUALIZATION AND ASSESSMENT

Visualizations of the importance of each region in an artwork can be generated using the regions of importance detection method described in the section “Visualization of Artist-Characteristic Regions,” where the occlusions are performed with a gray block of 8×8 pixels to indicate approximate regions that are characteristic of the artist. The regions of importance can be visualized using

heat map color coding, as shown in Figure 3(b). The value of a region in the heat map corresponds to the certainty score of PigeoNET for the artwork with that region occluded. In other words, a region with a lower value is of greater importance in correctly attributing the artwork, with (dark) red regions being highly characteristic of the artist, and (dark) blue regions being the least characteristic.

THE GOAL OF AN ARTIST ATTRIBUTION TASK IS TO ATTRIBUTE AN UNSEEN ARTWORK TO THE ARTIST WHO CREATED IT.

When comparing the artwork and heat map in Figure 3 of the drawing by Rembrandt, it is very noticeable that PigeoNET assigns much weight to seemingly empty areas. The texture of the material on which an artwork is created can be indicative of the artist who created the artwork

[26]. When taking a closer look at Figure 4, with enhanced contrast, it becomes apparent that the areas are not empty and that a distinctive visual texture is present. The visual pattern is sufficiently distinctive and artist-specific for PigeoNET to assign it a larger weight. The pattern is an example of a visual characteristic that is indirectly related to the artist. It illustrates the importance of the transparency of automatic attribution to allow human experts to interpret and evaluate the visual characteristic.

DECIDING BETWEEN TWO ARTISTS


In the “Visualization and Assessment” section, we used PigeoNET to attribute an artwork to a single artist. Yet, as illustrated by the work of Peter Paul Rubens and Jan Brueghel in Figure 1, in many cases two (or more) artists have worked on the same artwork (see also [27]).

As evident from our results, PigeoNET had difficulty in correctly attributing artworks of closely collaborating artists. An intriguing explanation for PigeoNET’s failure to assign the “correct one” of two potential artists to artworks is that the artworks are created by both artists. In that case, it would not be a failure at all and indicates that PigeoNET discovered that the two artists are similar, and it recognizes the characteristic features of both artists, even if the work is attributed to only one artist. In the remainder of this section, we demonstrate the possibility of using PigeoNET to perform a fine-grained analysis of an artwork, attributing individual image regions to an artist.

DISCOVERING DUAL AUTHORSHIP

PigeoNET had difficulty in distinguishing between the works of Jan and Caspar Luyken, a father and son who worked together and created many prints. Throughout their careers, Jan Luyken chose to depict pious and biblical subjects, whereas Caspar Luyken mostly depicted worldly scenes [28]. As an example, we consider the artwork shown in Figure 5, *Over dracht der Nederlande, aan de Infante Isabella*. The work depicts the transfer of the Spanish Netherlands by Filips II to Isabella Clara Eugenia. Although arguably it is a very worldly scene, it is nevertheless attributed to Jan Luyken. Could it be possible that the artwork is incorrectly attributed to Jan Luyken? Obviously, this is a question that has to be answered by experts of their works.

Our findings may support them in their assessment. Although PigeoNET correctly attributed the artwork to Jan Luyken, the reported certainty score for Caspar Luyken is very high. Apparently, PigeoNET responds to visual features that are characteristic of Caspar Luyken. Using PigeoNET’s visualization, we are able to determine for each region how characteristic it is for each of the two artists. We created a visualization based on the certainty scores for



[FIG4] A contrast-enhanced detailed view of a highly textured region of the artwork shown in Figure 3(a).



[FIG5] An image of the *Overdracht der Nederlande, aan de Infante Isabella* (1697–1699) by Jan Luyken.

Jan Luyken and Caspar Luyken. Figure 6 shows the visualization using color coding on a yellow to blue scale. The yellow regions are characteristic for Jan Luyken, whereas the blue regions are characteristic for Caspar Luyken, the green regions are indeterminate and show characteristics of either artists in equal amounts. This example demonstrates the potential use of PigeoNET to support the study of dual authorship artworks.

DISCUSSION

Previous work on automatic artist attribution has shown that prior knowledge can be leveraged to engineer features for automatic artist attribution. In this article, we presented a novel approach that does not rely on prior knowledge and is capable of discovering characteristic features automatically enabling a successful artist attribution. Additionally, we demonstrated that PigeoNET visualizations reveal artwork regions most characteristic of the artist and that PigeoNET can aid in answering outstanding questions regarding dual-authorship.

We will now discuss the considerations regarding the data set used and address how the selection of subsets may affect the nature of visual characteristics discovered. Although the Rijksmuseum Challenge data set is the largest available data set containing digital reproductions of artworks acquired under controlled conditions [24], it does suffer from two main limitations. First, given the wide variety of artwork types, it is unclear how the “controlled conditions” were defined for different artworks. Any variation in the reproduction setting (e.g., illumination, perspective, camera type) may be picked up by PigeoNET. Presumably, our P (prints only) data sets suffer less from this problem. Still, even in these data sets subtle differences in

**THE RESULTS OBTAINED
IN THIS WORK ON THE
AUTOMATIC ARTIST ATTRIBUTION
TASK SHOW THAT PIGEONET
IS CAPABLE OF ACCURATELY
ATTRIBUTING UNSEEN WORKS
TO THE CORRECT ARTIST.**


proper visualizations are essential to aid art experts in their assessment of the feasibility of classifications.

The second limitation concerns the labeling of artworks. After having performed our main experiments, we discovered that for some artworks, the Rijksmuseum catalog lists multiple contributions, whereas the Rijksmuseum Challenge data set only lists a single artist [24].

The contributions listed in the Rijksmuseum catalog vary greatly (from inspiration to dual authorship) and do not always influence the actual attribution, but do create uncertainty about the attribution of artworks in the Rijksmuseum challenge data set. Although this significantly limits the possibility of learning stylistic features from such artworks, it does not prohibit PigeoNET from learning visual characteristics that are associated with the primary artist as such characteristics remain present in the artwork. Still, the validity and consistency of attributions is of major concern to safeguard the validity of methods such as PigeoNET. Also in the creation of such databases, the involvement of human art experts is required.

The results obtained in this work on the automatic artist attribution task show that PigeoNET is capable of accurately attributing unseen works to the correct artist. The increase of performance for the sets with a higher number of examples shows that including more examples per artist leads to a better performance. Moreover, the complete Rijksmuseum Challenge data set is a highly diverse data set with many different types of art. For some cases (e.g., the porcelain of the Meissener Porzellan Manufaktur) this results in a class that is visually very distinctive from the rest of the data set, which could make it easier to identify the correct artist. However, when comparing the performances obtained on the homogeneous P1 subsets (prints only) with those on the more heterogeneous A subsets (all artwork types), the difference in performance is quite small. This demonstrates that PigeoNET is capable of learning a rich representation of multiple artwork types without a major impact on its predictive power. Part of the types of features discovered in the A subsets are likely to distinguish between art types (e.g., a porcelain object versus a painting), rather than between author styles. In the P subsets, features will be more tuned to stylistic differences, because these subsets are confined to a single type of artwork.

Our findings indicate that the number of artists and the number of examples per artist have a very strong influence on the performance, which suggests that a further improvement of the performance is possible by expanding the data set. In future research we will determine to what extent this is the case.



[FIG6] Visualization of how characteristic each image region is for the artists Jan and Caspar Luyken. The yellow regions are characteristic of Jan Luyken, whereas the blue regions are characteristic of Caspar Luyken.

CONCLUSIONS

In this article we have evaluated a feature learning system to assess to what extent it is possible to discover an artist's visually characteristic features. The results on the automatic attribution task demonstrate that the system is capable, up to a high degree of accuracy, of using visual characteristics to assign unseen artworks to the correct artist. Moreover, we demonstrated the possibility of using the visual characteristics to reveal the artist of a specific region within an artwork, which in the case of multiple artists could lead to new discoveries about the origin and creation of important works of cultural heritage. Therefore, we conclude that PigeoNET represents a fruitful approach for future computer-supported examination of artworks.

ACKNOWLEDGMENTS

We thank Teio Meedendorp and Louis van Tilborgh for their useful contributions to earlier versions of this article. Furthermore, we acknowledge the anonymous reviewers for their helpful comments. The research reported in this article is performed as part of the Reassessing Vincent van Gogh project supported by The Netherlands Organisation for Scientific Research (NWO; grant 323.54.004) in the context of the Science4Arts research program.

AUTHORS

Nanne van Noord (n.j.e.vannoord@tilburguniversity.edu) is a Ph.D. candidate working on the Reassessing Vincent van Gogh project at the Tilburg Center for Cognition and Communication, Tilburg University, The Netherlands. His research interests include machine learning, representation learning, and computer vision.

Ella Hendriks (hendriks@vangoghmuseum.nl) is a senior paintings conservator at the Van Gogh Museum, as well as an associate professor at the University of Amsterdam Training Programme for Conservation and Restoration of Cultural Heritage.

Eric Postma (e.o.postma@tilburguniversity.edu) is a full professor in artificial intelligence at the Tilburg Center for Cognition and Communication, Tilburg University, The Netherlands. His work focuses on machine learning and image processing with an emphasis on cultural heritage applications.

REFERENCES

- [1] J. Rush, "Acquiring a concept of painting style," Ph.D. dissertation, Dept. of Educational Psychology, Univ. Arizona, Tucson, AZ, 1974.
- [2] I. Berezchnoy, E. Postma, and J. van den Herik, "Computer analysis of Van Gogh's complementary colours," *Pattern Recog. Lett.*, vol. 28, no. 6, pp. 703–709, Apr. 2007.
- [3] J. Li, L. Yao, E. Hendriks, and J. Z. Wang, "Rhythmic brushstrokes distinguish van Gogh from his contemporaries: findings via automated brushstroke extraction," *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 34, no. 6, pp. 1159–76, June 2012.
- [4] M. M. van Dantzig, *Pictology. An Analytical Method for Attribution and Evaluation of Pictures*, the van Dantzig Foundation, editor. Leiden, The Netherlands: Brill Leiden, 1973.
- [5] T. Lombardi, "The classification of style in fine-art painting," Ph.D. dissertation, School of Computer Science and Information Systems, Pace Univ., NY, 2005.
- [6] C. R. Johnson, Jr., E. Hendriks, I. J. Berezchnoy, E. Brevdo, S. M. Hughes, I. Daubechies, J. Li, E. Postma, and J. Z. Wang, "Image processing for artist identification," *IEEE Signal Processing. Mag.*, vol. 25, no. 4, pp. 37–48, 2008.
- [7] J. M. Hughes, D. J. Graham, and D. N. Rockmore, "Quantification of artistic style through sparse coding analysis in the drawings of Pieter Bruegel the Elder," *Proc. Nat. Acad. Sci. U.S.A.*, vol. 107, no. 4, pp. 1279–83, Jan. 2010.
- [8] R. Taylor, R. Guzman, and T. Martin, "Authenticating Pollock paintings using fractal geometry," *Pattern Recog. Lett.*, vol. 28, no. 6, pp. 695–702, 2007.
- [9] L. J. P. van der Maaten and E. O. Postma, "Texon-based analysis of paintings," in *SPIE Optical Imaging and Applications*, Aug. 2010, vol. 7798. [Online]. Available: <http://proceedings.spiedigitallibrary.org/volume.aspx?conferenceid=1936&volumeid=678>
- [10] P. Abry, H. Wendt, and S. Jaffard, "When Van Gogh meets Mandelbrot: Multifractal classification of painting's texture," *Signal Process.*, vol. 93, no. 3, pp. 554–572, Mar. 2013.
- [11] M. Erait, "Non-formal learning and tacit knowledge in professional work," *Br. J. Educ. Psychol.*, vol. 70, no. 1, pp. 113–136, Mar. 2000.
- [12] Y. Bengio, A. Courville, and P. Vincent, "Representation learning: A review and new perspectives," *IEEE Trans. Pattern Anal. Machine Intell.*, vol. 35, no. 8, pp. 1798–1828, Aug. 2013. doi:10.1109/TPAMI.2013.50
- [13] A. Krizhevsky, I. Sutskever, and G. Hinton, "ImageNet classification with deep convolutional neural networks," in *Adv. Neural Inform. Process. Syst.*, vol. 25, 2012, pp. 1097–1105.
- [14] S. Watanabe, J. Sakamoto, and M. Wakita, "Pigeons' discrimination of paintings by Monet and Picasso," *J. Exp. Anal. Behav.*, vol. 63, no. 2, pp. 165–174, 1995.
- [15] W. Wu, A. M. Moreno, J. M. Tangen, and J. Reinhard, "Honeybees can discriminate between Monet and Picasso paintings," *J. Comparative Physiol. A*, vol. 199, no. 1, pp. 45–55, Oct. 2012.
- [16] M. Zeiler and R. Fergus, "Visualizing and understanding convolutional neural networks," 2013, arXiv:1311.2901v2.
- [17] Y. Bengio and Y. Lecun, "Convolutional Networks for Images, Speech, and Time-Series," in *The Handbook of Brain Theory and Neural Networks*, M. A. Arbib, Ed. Cambridge, MA: MIT Press, 1995, pp. 255–258.
- [18] Y. LeCun, B. Boser, J. S. Denker, D. Henderson, R. E. Howard, W. Hubbard, and L. D. Jackel, "Backpropagation applied to handwritten zip code recognition," *Neural Computat.*, vol. 1, no. 4, pp. 541–551, Dec. 1989.
- [19] A. K. Jain, N. K. Ratha, and S. Lakshmanan, "Object detection using Gabor filters," *Pattern Recog.*, vol. 30, no. 2, pp. 295–309, Feb. 1997.
- [20] M. Oquab, L. Bottou, I. Laptev, and J. Sivic, "Learning and transferring mid-level image representations using convolutional neural networks," in *Proc. 2014 IEEE Conf. Computer Vision Pattern Recognition*, June 2014, pp. 1717–1724.
- [21] K. Simonyan, A. Vedaldi, and A. Zisserman, "Deep inside convolutional networks: visualising image classification models and saliency maps," pp. 1–8, Dec. 2013, arXiv:1312.6034.
- [22] D. Erhan and Y. Bengio, "Visualizing higher-layer features of a deep network," Tech. Rep. 1341, Univ. Montreal, 2009.
- [23] A. T. Woollett and A. van Suchtelen, *Rubens and Brueghel: A Working Friendship*. Los Angeles, CA: J. Paul Getty Museum, 2006.
- [24] T. Mensink and J. van Gemert, "The Rijksmuseum challenge: museum-centered visual recognition," in *Proc. Int. Conf. Multimedia Retrieval*, 2014, pp. 2–5.
- [25] Y. Jia, (2013). Caffe: An open source convolutional architecture for fast feature embedding. [Online]. Available: <http://caffe.berkeleyvision.org>
- [26] D. H. Johnson, C. R. Johnson, Jr., A. G. Klein, W. A. Sethares, H. Lee, and E. Hendriks, "A thread counting algorithm for art forensics," in *Proc. IEEE 13th Digital Signal Processing Workshop and 5th IEEE Signal Processing Education Workshop*, 2009, pp. 679–684.
- [27] H. Bevers, L. Hendrix, W. W. Robinson, and P. Schatborn, *Drawings by Rembrandt and His Pupils: Telling the Difference*. Los Angeles, CA: J. Paul Getty Museum, 2009.
- [28] [Online]. Available: http://www.geheugenvannederland.nl/?en/collecties/jan_en_casper_luyken/vader_zoom